DA5402 - Assignment 6

Anirudh Rao be21b004

Task 1

To scrape the required IO and CPU metrics, a function called `collect_iostat_metrics()` was created. This uses the `subprocess` library to run `iostat`. The output of this is parsed to identify the specific metrics. The five IO metrics are recorded for each device. In the case of CPU metrics, the average percentage is found for each mode except for the '`steal`' mode, as mentioned in the assignment. The scraped metrics are assigned to Prometheus gauges that are exposed at http://localhost:18000/ using a Prometheus client. This scraping is repeated every 1 second using a thread (from the `threading` library). Logging is implemented to record successful scraping or errors that occur during scraping.

Task 2

Memory information metrics are collected using a function called `collect_meminfo_metrics()`. This scrapes the metrics by reading and parsing the `/proc/meminfo` file into key-value pairs. The keys (e.g., `MemFree`, `Active(anon)`, `SwapTotal`) are normalized into Prometheus-friendly metric names that have the prefix '`meminfo_`'. This is done using a helper function called `normalize_meminfo_key()`. The scraped metrics are assigned to Prometheus gauges that are exposed at http://localhost:18000/ using a Prometheus client. This scraping is repeated every 1 second using a thread. Logging is implemented to record successful scraping or errors that occur during scraping.

Task 3

To setup a Prometheus server, a `prometheus.yml` file was created. This specifies the scrape interval as 2 seconds and the scraping target as http://localhost:18000/. The metrics exposed by the Python script can be queried at the Prometheus UI server. This server is available on http://localhost:9090/ after Prometheus is run using the provided `prometheus.yml` file.

Files submitted

- `requirements.txt` – Text file that contains the Python libraries to be installed prior to running the script.
- `main.py` – Python script that contains the functions involved in scraping the metrics from the Linux environment. Exposes the scraped metrics at http://localhost:18000/ using a Prometheus client. Creates a logfile called `script.log` to collect logging information.
- `prometheus.yml` – YML file that contains the configuration for setting up the Prometheus server with the required scraping target and frequency. This server will be available on http://localhost:9090/.

Steps to run code

- Open a Linux environment.
- Ensure that Python3, `pip`, and Prometheus are installed in the Linux environment.
- Install the requirements given in `requirements.txt` using `pip`.
- Run `main.py`.
- Verify that `main.py` is running by checking http://localhost:18000/ or the `script.log` logfile.
- In the Linux environment, run Prometheus using the command `./prometheus --config.file=./prometheus.yml`.
- Check the Prometheus UI server at http://localhost:9090/ and query for the relevant metrics.
- To stop the collection of metrics, press Ctrl. + C for both the `main.py` script and the Prometheus server.

Sample outputs

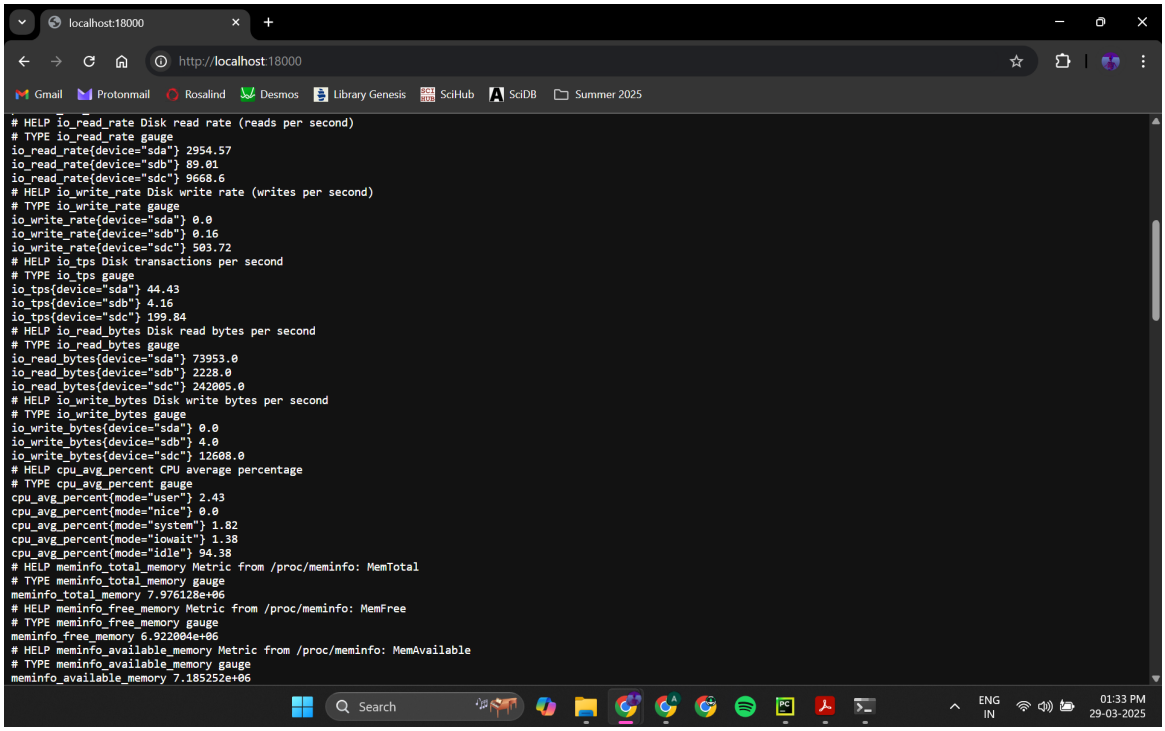

Figure 1: The scraped metrics exposed at http://localhost:18000/

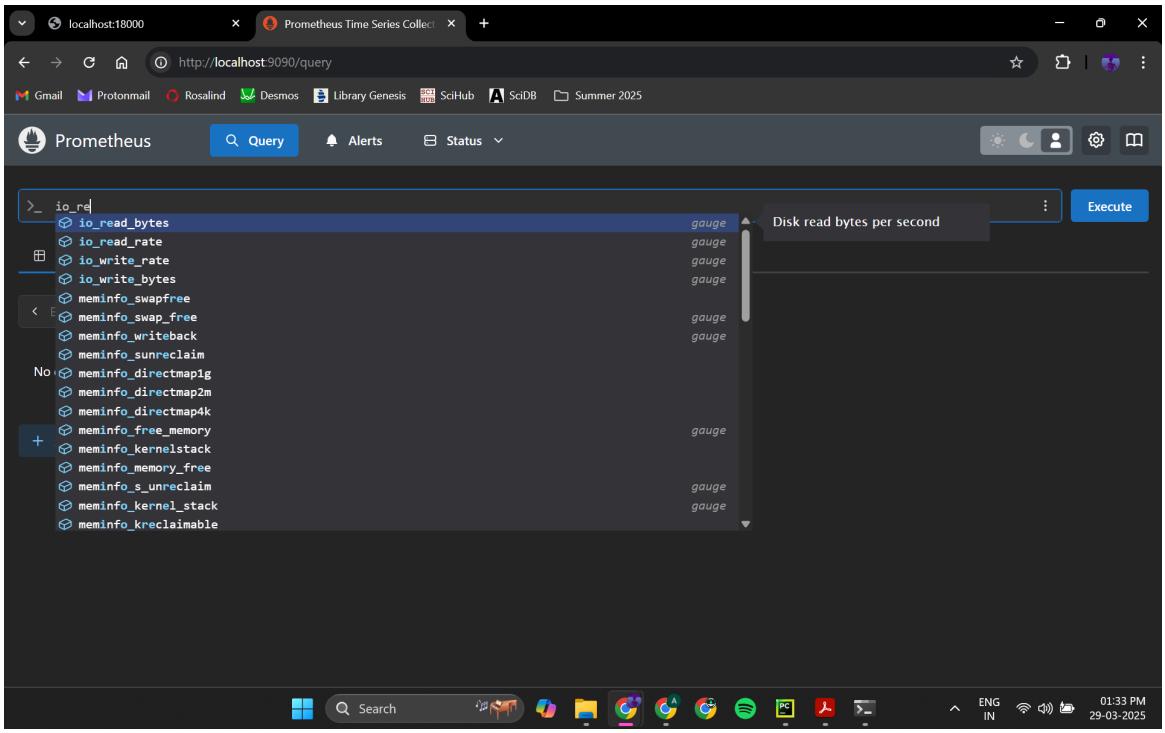

Figure 2: Querying the metrics in the Prometheus UI server at http://localhost:9090/

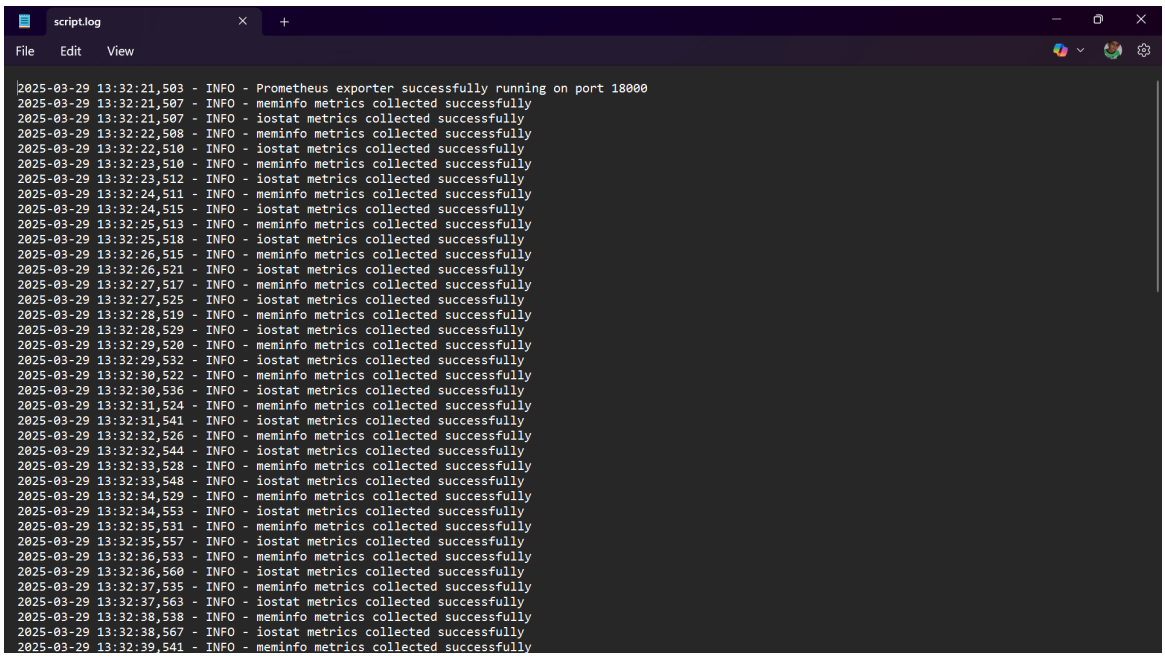

Figure 3: The `script.log` file capturing the logging information

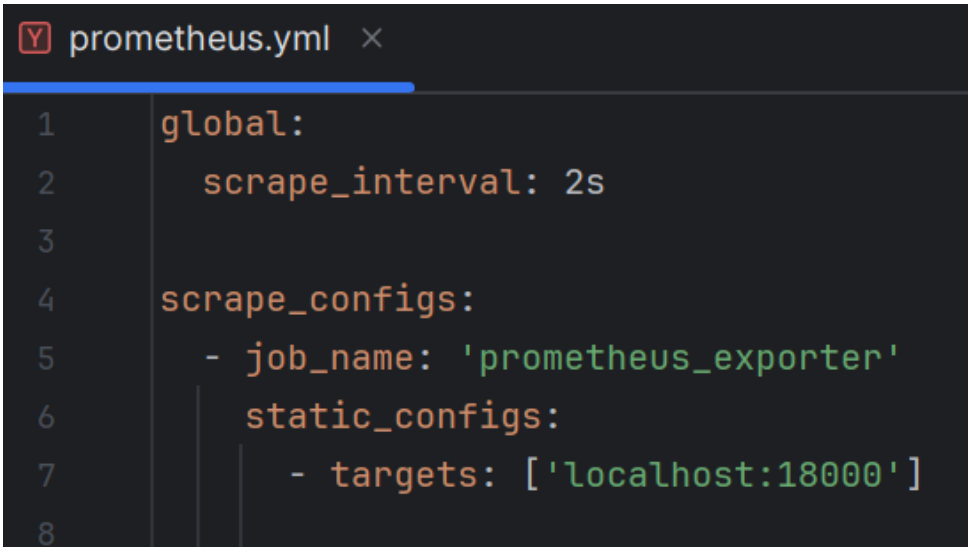

Figure 4: The `prometheus.yml` file containing the configuration for the Prometheus server